

---

# Data Structures and Algorithms in Java™

Sixth Edition

**Michael T. Goodrich**

Department of Computer Science  
University of California, Irvine

**Roberto Tamassia**

Department of Computer Science  
Brown University

**Michael H. Goldwasser**

Department of Mathematics and Computer Science  
Saint Louis University

---

## Instructor's Solutions Manual

WILEY

## Chapter

# 9

## Maps, Hash Tables, and Skip Lists

### Hints and Solutions

#### Reinforcement

**R-9.1) Hint** The first insertion is  $O(1)$ , the second is  $O(2)$ , ...

**R-9.2) Hint** Use the `PositionalList.remove` method to delete an entry from the map.

**R-9.3) Hint** Take advantage of the existing `findIndex` method.

**R-9.3) Solution**

```
public boolean containsKey(K key) {  
    return (findIndex(key) != -1);  
}
```

**R-9.4) Hint** Think about which of the schemes use the array supporting the hash table exclusively and which of the schemes use additional storage external to the hash table.

**R-9.5) Hint** There is a lot of symmetry and repetition in this string, so avoid a hash code that would not deal with this.

**R-9.5) Solution** Either polynomial hash codes or cyclic shift hash codes would be good.

**R-9.6) Hint** Try to mimic the figure in the book.

**R-9.6) Solution**

0	1	2	3	4	5	6	7	8	9	10
13	94				44			12	16	20
	39				88			23	5	
					11					

**R-9.7) Hint** Try to mimic the figure in the book.

**R-9.7) Solution**

0	1	2	3	4	5	6	7	8	9	10
13	94	39	16	5	44	88	11	12	23	20

**R-9.8) Hint** The failure occurs because no empty slot is found. For the drawing, try to mimic the figure in the book.

**R-9.8) Solution**

0	1	2	3	4	5	6	7	8	9	10
13	94	39	11		44	88	16	12	23	30

The probe sequence when inserting value 5 fails to ever find the sole remaining empty slot (at index 4).

**R-9.9) Hint** Try to mimic the figure in the book.

**R-9.9) Solution**

0	1	2	3	4	5	6	7	8	9	10
13	94	23	88	39	44	11	5	12	16	20

**R-9.10) Hint** Think of the worst-case time for inserting every entry in the same cell in the hash table.

**R-9.10) Solution** The worst-case time is  $O(n^2)$ . The best case is  $O(n)$ .

**R-9.11) Hint** Mimic the way the figure is drawn.

**R-9.11) Solution**

0	1	2	3	4	5	6	7	8	9	10	11	12	13	14	15	16	17	18
		12	18	41	30	36	25		54			38	10			28		

**R-9.12) Hint** Combine the hash values of the two components to make a hash value for the pair.

**R-9.12) Solution**

```
public int hashCode() {
    return (a.hashCode() ^ b.hashCode());
}

public boolean equals(Object o) {
    if (o == null) return false;
    if (getClass() != o.getClass()) return false;
    Pair other = (Pair) o;
    return (a.equals(other.a) && b.equals(other.b));
}
```

**R-9.13) Hint** There is a subtle flaw—but can you find it?

**R-9.13) Solution** The problem occurs if null is allowed as a value within the map. In that case, there is ambiguity when `bucket.remove(k)` returns null, as that may indicate that nothing was removed, or that a key was removed having null as its value. Therefore, we might need to decrease `n`, even when null is returned.

**R-9.14) Hint** The load factor can be controlled from within the abstract class, but there must be means for setting the parameter (either through the constructor, or a new method).

**R-9.15) Hint** It is okay to insert a new entry on “top” of the deactivated entry object.

**R-9.16) Hint** You will need to keep track of the number of probes in order to apply quadratic probing.

**R-9.17) Hint** Think of where the entry with minimum key is stored.

**R-9.18) Hint** Since the map will still contain  $n$  entries at the end, you can assume that each `remove()` operation takes the same asymptotic time.

**R-9.19) Hint** Take advantage of the existing `findIndex` method.

**R-9.19) Solution**

```
public boolean containsKey(K key) {
    int j = findIndex(key);
    return (j < table.size() && compare(key, table.get(j)) == 0);
}
```

**R-9.20) Hint** The crucial methods are `get` and `put`.

**R-9.21) Hint** In the new code, what happens when `key` equals `table.get(mid)`?

**R-9.21) Solution** This version is certainly wrong. This can be seen by considering for example a call of `findIndex(20, 0, 2)` for a table with contents `{10, 20, 30}`. Such a call should return index 1, yet the newly proposed code will return index 3.

It is worth noting that the alternate version would produce the correct result if only line 4 had been written as

```
if (compare(key, table.get(mid)) <= 0)
```

thereby grouping the equal case with the `if` clause (not the `else` clause).

**R-9.22) Hint** Assume that a skip list is used to implement the sorted map.

**R-9.22) Solution** For this problem, we assume we are using a skip list as the implementation for the sorted map. (If a `SortedTableMap` were used, with lower price/performance pairs, each pair is inserted at the beginning of the array list, and thus the insertion time dominates the search time, for overall  $O(n^2)$  time.)

With a skip list, the expected running time for maintaining a maxima set if we insert  $n$  pairs such that each pair has lower price and performance than one before it is  $O(n \log n)$ , since each pair is inserted and does not remove any others. Thus, the  $i^{\text{th}}$  insertion runs in  $O(\log i)$  time, and the total expected running time is  $\sum_{i=1}^n \log i$ , which is  $O(n \log n)$ .

If, on the other hand, each pair had a lower price and higher performance than the one before it, then the overall running time is  $O(n)$ , because each new pair causes the removal of the pair that came before it. Therefore the size of the map is constant, and thus all its operations run in expected  $O(1)$  time.

**R-9.23) Hint** Mimic the style of the figures in the book.

**R-9.24) Hint** You must link out the removed entry's tower from all the lists it belongs to.

**R-9.25) Hint** Compare to the implementation of the `addAll` method given in Section 10.5.1.

**R-9.25) Solution** Loop through the second set, while removing each of its elements from the first set (if it exists).

**R-9.26) Hint** Compare to the implementation of the `addAll` method given in Section 10.5.1.

**R-9.26) Solution** Loop through the first set, and remove each element that cannot be found in the second set.

**R-9.27) Hint** Recall that most skip-list operations run in  $O(\log k)$  time for a set with size  $k$ .

**R-9.27) Solution**  $O(m \log(n + m))$  time.

**R-9.28) Hint** Recall that most hash-table operations run in  $O(1)$  expected time.

**R-9.28) Solution**  $O(m)$

**R-9.29) Hint** Recall that most hash-table operations run in  $O(1)$  expected time.

**R-9.29) Solution**  $O(m)$

**R-9.30) Hint** Recall that most hash-table operations run in  $O(1)$  expected time.

**R-9.30) Solution**  $O(n)$

**R-9.31) Hint** Something from this chapter should be helpful!

**R-9.31) Solution** We should use a sorted multiset, as we want to maintain multiple entries with the same birthday and we want to be able to do neighboring queries.

---

## Creativity

**C-9.32) Hint** Consider a bootstrapping method for finding the primes up to  $\sqrt{2M}$ .

**C-9.33) Hint** Your solution should make a single call to `findIndex`.

**C-9.33) Solution**

```

public V putIfAbsent(K key, V value) {
    int j = findIndex(key);
    if (j == -1) {
        table.add(new MapEntry<>(key, value));    // add new entry
        return null;
    } else {
        return table.get(j).getValue();           // return existing value
    }
}

```

**C-9.34) Hint** You may assume that such a method is supported by the auxiliary `UnsortedTableMap`

**C-9.35) Hint** You must only call the `findSlot` utility once.

**C-9.36) Hint** Model your solution after the existing support for other map methods.

**C-9.37) Hint** You may assume that such a method is supported by the auxiliary `UnsortedTableMap`

**C-9.38) Hint** The heart of the process can be performed by the `findSlot` utility.

**C-9.39) Hint** When might the load factor fall below the threshold, and how can you detect this from within `AbstractHashMap`?

**C-9.40) Hint** Start by defining an appropriate subclass of `AbstractMap.MapEntry` that includes the new `next` field that is desired.

**C-9.41) Hint** You need to do some shifting of entries to close up the “gap” just made, but you should only do this for entries that need to move.

**C-9.41) Solution** When we remove an entry from a hash table that uses linear probing without using deleted element markers, we must ensure that any cluster of consecutively filled cells remains intact, unless no entries in that cluster after a newly introduced gap have a hash value that falls before such a gap. Assume that we wish to delete an entry stored at index  $j$  and that the cluster it belongs to goes until an empty cell at index  $k$ . For ease of exposition only, we assume  $k > j$ , thus the cluster does not wrap around the end of the array. To fill the hole in index  $j$ , we wish to find the next entry of the cluster that has a hash value  $h \leq j$ . If no such entry exists, we may leave a hole at  $j$ . Otherwise, assume index  $i$  holds the first such entry. We move that entry to  $j$  and then we repeat the process to fill the hole that results at index  $i$ , finding the next subsequent entry of the cluster with hash value  $h \leq i$ . Note that the entire process takes time linear in the length of the original cluster of filled cells.

**C-9.42) Hint** For part a, note that the symmetry will halve the range of possible values. For part b, note that such automatic collisions will not occur.

**C-9.43) Hint** Perhaps borrow techniques from the Progression hierarchy of Chapter 2, or use Java's iterators.

**C-9.44) Hint** Maintain a secondary PositionalList instance that represents the FIFO order, and store positions within that list with each entry of the hash table.

**C-9.45) Hint** Each map entry instance should store its current index within the table.

**C-9.46) Hint** Each map entry instance should store its current index within the table.

**C-9.47) Hint** Each map entry instance should store its position within the bucket list.

**C-9.48) Hint** Try out some examples.

**C-9.48) Solution** The original version of Code Fragment 10.11 does not guarantee that it finds the leftmost occurrence when duplicates exist, because it stops searching as soon as it finds a match. For example, if the entire table was filled with duplicates, it will report the middle index.

If the version in Exercise R-9.21 had been written as intended, with the  $\leq$  operator on line 4, then it is a correct implementation and indeed guarantees that it finds the leftmost of all duplicates.

**C-9.49) Hint** Do a "double" binary search.

**C-9.49) Solution** In order to find the  $k^{\text{th}}$  smallest key in the union of the keys from  $S$  and  $T$ , we can do a "double" binary search. In other words, we will begin by examining the  $k/2^{\text{th}}$  element in the array list  $S$ . Next, we will find the largest element in  $T$  that is less than  $S[k/2]$  by binary search. Then, we will add the indices of the elements we were examining in  $S$  and  $T$ . If the sum equals  $k$ , then the max of the two elements is our result. If the sum is greater than  $k$ , then we will do a binary search to the right (upwards) in  $S$ . If the sum is less than  $k$ , then we will do a binary search to the left (downwards) in  $S$ . This is followed once again by searching in  $T$  for largest element less than the current element in  $S$ , etc. This method does a binary search on  $S$  which requires  $O(\log n)$  "probes." However, for each probe of the search, it does a binary search on  $T$  which takes  $O(\log n)$  time. Thus, the entire method takes  $O(\log^2 n)$  time.

**C-9.50) Hint** Dove-tail two binary searches.

**C-9.51) Hint** Do a lazy iteration through indices of the underlying array list.

**C-9.52) Hint** Manage a primary iteration through all nonempty buckets, and then a secondary iteration through each element of a bucket.

**C-9.53) Hint** Do a lazy iteration through the nonempty cells of the table.

**C-9.54) Hint** Think about first sorting the pairs by cost.

**C-9.54) Solution** Sort the pairs by cost. Then scan this list looking at the performance values. Remove any that have performance values worse than the (unremoved) pair that came before.

**C-9.55) Hint** In the insertion algorithm, first repeatedly flip the coin to determine the level where you should start inserting the new entry.

**C-9.56) Hint** Consider augmenting each node  $v$  in a higher level with the number of missing entries in the gap from  $v$  to the next node over.

**C-9.57) Hint** Consider augmenting each node  $v$  in a higher level with the number of missing entries in the gap from  $v$  to the next node over.

**C-9.58) Hint** Think first about how you can determine the number of 1's in any row in  $O(\log n)$  time.

**C-9.58) Solution** To count the number of 1's in  $A$ , we can do a binary search on each row of  $A$  to determine the position of the last 1 in that row. Then we can simply sum up these values to obtain the total number of 1's in  $A$ . This takes  $O(\log n)$  time to find the last 1 in each row. Done for each of the  $n$  rows, then this takes  $O(n \log n)$  time.

**C-9.59) Hint** Consider a two-pass solution, with use of an auxiliary structure.

**C-9.59) Solution**

```
public void retainAll(Set<E> other) {
    // let's first build a list of elements to remove
    ArrayList<E> leaving = new ArrayList<E>();
    for (E element : this)
        if (!other.contains(element))
            leaving.add(element);
    // now let's remove those elements from the set
    for (E element : leaving)
        remove(element);
}
```

**C-9.60) Hint** Think of describing your algorithms in terms of boolean operations on the bit vectors.

**C-9.61) Hint** Think of how you could transform  $D$  into  $L$ .

**C-9.62) Hint** Consider the way `subMap` was implemented for `SortedTableMap`.

**C-9.62) Solution** Simply do a binary search to find an element equal to  $k$ . Then step back through the array until you reach the first element equal to



$k$ . Finally, step forward through the area reporting entries until you reach the first key that is not equal to  $k$ . This takes  $O(\log n)$  time for the search and then at most  $s$  time to search back to the beginning of the run of  $k$ 's and  $s$  time return all of the elements with key  $k$ . Therefore, we have a solution running in at most  $O(s + \log n)$  time.

**C-9.63) Hint** You need some way of grouping together entries with the same key.

**C-9.63) Solution** Use a skip list map, with only  $r$  nodes in the bottom level, and with each of these having a secondary linked list of elements with the same key. Thus, all the query and update operations will run in  $O(\log r)$  expected time, and the `getAll` method will run in  $O(\log r + s)$  expected time.

**C-9.64) Hint** You need some way of grouping together elements with the same key.

**C-9.64) Solution** The solution is actually quite simple—just use something like the `HashMultimap` class given in the book, but make sure that the capacity of the underlying hash table is always  $O(n)$ . Another possibility is to use a linked hash table, which keeps its entries in a linked list in addition to the hash table. Either way we would get that the `entries()` method would run in  $O(n)$  time.

---

## Projects

**P-9.65) Hint** The biggest challenge is detecting the case of an infinite loop.

**P-9.66) Hint** When searching for an existing key, make sure to consider both of the possible buckets.

**P-9.67) Hint** Maintain a secondary `PositionalList` instance that represents the FIFO order, and store positions within that list with each entry of the hash table.

**P-9.68) Hint** We've already implemented them for you; all that's left is the experiments.

**P-9.69) Hint** In a Unix/Linux system, a good place to start is `/usr/dict`.

**P-9.70) Hint** It is okay to generate these phone numbers more-or-less at random.

**P-9.71) Hint** Sentinels can be used in place of the theoretical  $-\infty$  and  $+\infty$ .

**P-9.72) Hint** Try to make your screen images mimic the skip list figures in the book.

**P-9.73) Hint** You need to find some way to let the intermediate nodes in the skip list keep track of the number of elements that have been skipped over.

**P-9.74) Hint** For each word  $t$  that results from a minor change to  $s$ , you can test if  $t$  is in  $W$  in  $O(1)$  time.