

Data Structures and Algorithms

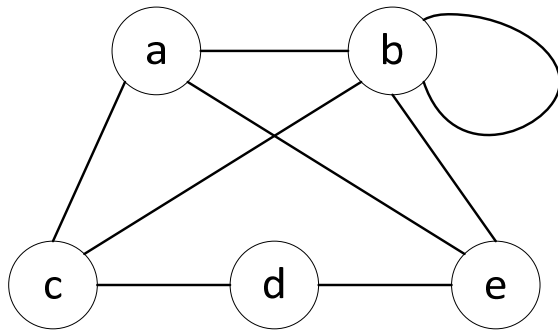
Chapter 14

Graph Algorithms

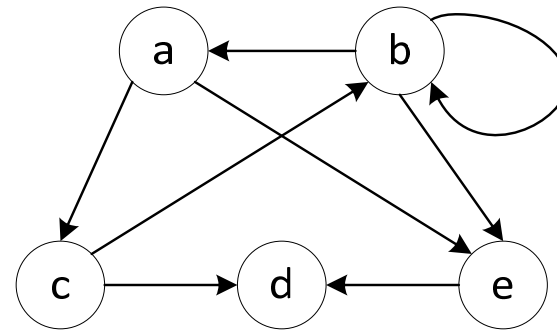
Basics

- A *graph* is a set V of *vertices* and a collection E of *edges*, $G = (V, E)$
- An edge connecting vertices (or nodes) u and v is denoted (u, v) .
- An edge can be *directed* or *undirected*.
- Directed graph vs. undirected graph:

(a) Undirected graph




(b) Directed graph



Graph Algorithms

Basics

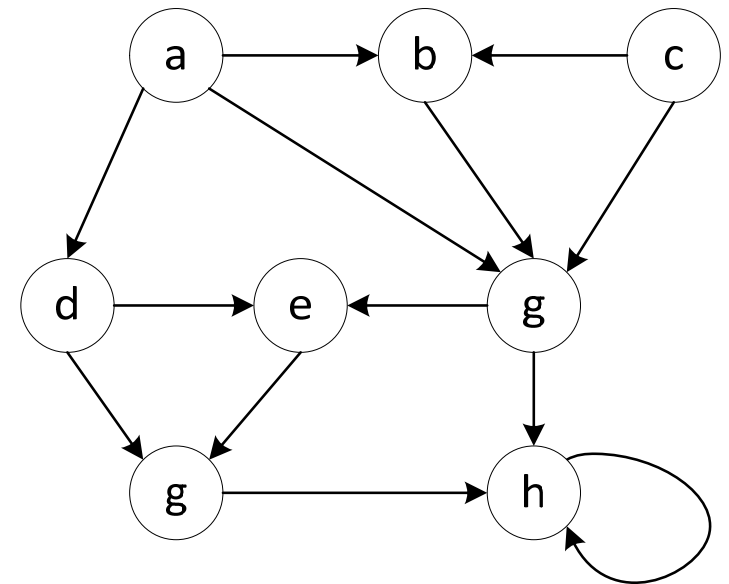
- Two vertices u and v are said to be *adjacent* if there is an edge (u, v) .
- An edge is said to be *incident* to a vertex if the vertex is one of the edge's endpoints.
- *Weighted* graph: An information (usually called weight) is associated with edges



Graph Algorithms

Basics


- Outgoing edge vs. incoming edge
 - Degree, in-degree, and out-degree of a node
-
- The outgoing edges of vertex g are (g, e) , (g, h) .
 - The incoming edges of vertex g are (a, g) , (b, g) , (c, g) .
 - The degree of vertex g , $\deg(g) = 5$.
 - The in-degree of vertex g , $\text{indeg}(g) = 3$.
 - The out-degree of vertex g , $\text{outdeg}(g) = 2$.



Graph Algorithms

Basics

- Parallel edges and self-loops

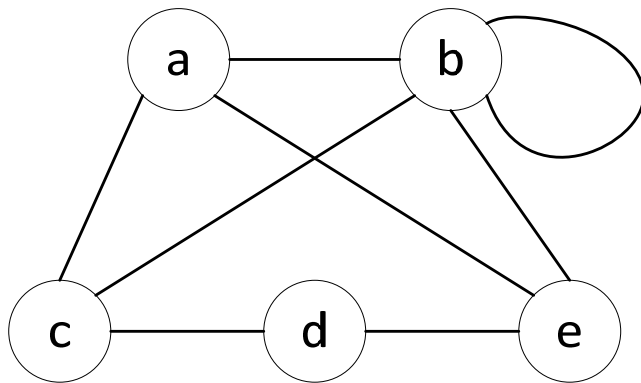


Graph Algorithms

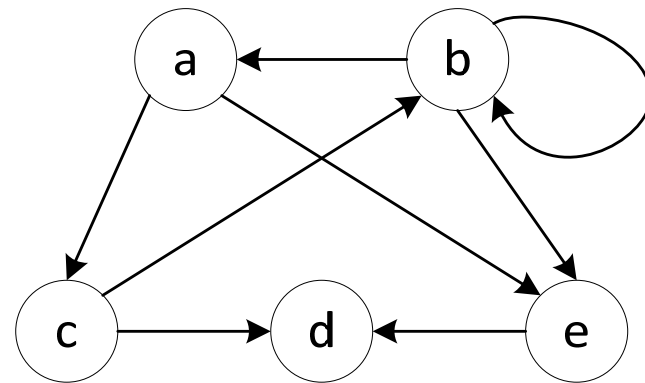
Basics

- Path, cycle, simple path, simple cycle, directed path, directed cycle*

(a) Undirected graph



(b) Directed graph




Graph Algorithms


Basics

- *Connected graph and strongly connected graph*

(a) Connected graph



(b) Strongly connected graph




Graph Algorithms


Basics

- *Subgraph and spanning subgraph*

(a) A subgraph



(b) A spanning subgraph




Graph Algorithms


Basics

- *Forest and tree*

(a) A forest



(b) A tree



- A *spanning tree* of a graph is a spanning subgraph that is a tree

Graph Algorithms

Basics

- Graph properties
 - If a graph $G = (V, E)$ has m edges, then
$$\sum_{v \in V} \deg(v) = 2m$$
 - If $G = (V, E)$ is a directed graph with m edges, then
$$\sum_{v \in V} \text{in deg}(v) = \sum_{v \in V} \text{out deg}(v) = m$$
- Let G be a simple graph with n vertices and m edges.
 - If G is undirected, then $m \leq \frac{n(n-1)}{2}$
 - If G is directed, then $m \leq n(n-1)$.

Graph Algorithms

Basics

- Graph properties (continued)
 - Let G be an undirected graph with n vertices and m edges:
 - If G is connected, then $m \geq n - 1$
 - If G is a tree, then $m = n - 1$
 - If G is a forest, then $m \leq n - 1$

Graph Algorithms

Graph ADT

- Operations
 - numVertices()
 - vertices()
 - numEdges()
 - edges()
 - getEdge(u , v)
 - endVertices(e)
 - opposite(v , e)

Graph Algorithms

Graph ADT

- Operations (continued)
 - `outDegree(v)`
 - `indegree(v)`
 - `outgoingEdges(v)`
 - `incomingEdges(v)`
 - `insertVertex(x)`
 - `insertEdge(u, v, x)`
 - `removeVertex(v)`
 - `removeEdge(e)`

Graph Algorithms

Data Structures for Graphs

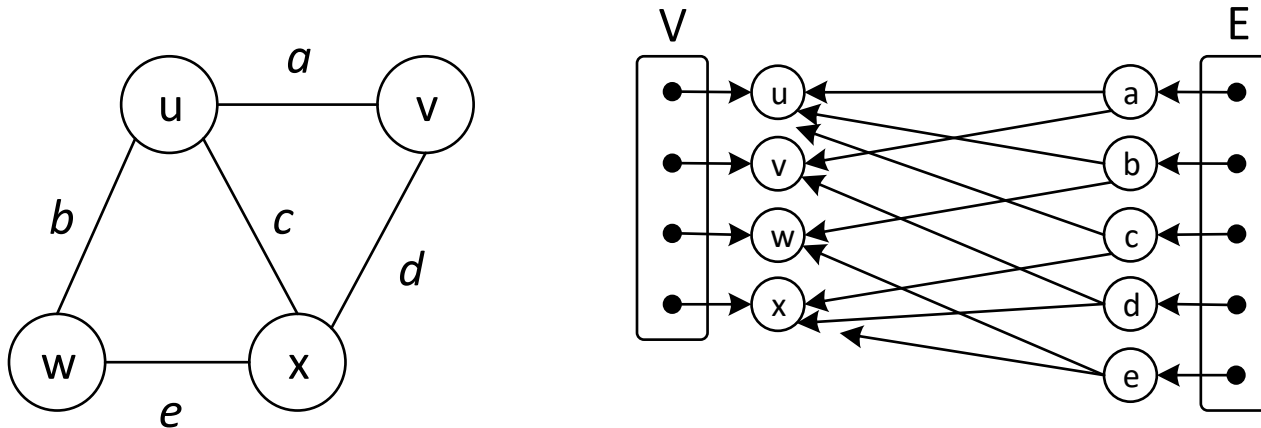
- Edge list, adjacency list, adjacency map, adjacency matrix

Method	Edge List	Adj. List	Adj. Map	Adj. Matrix
vertices()	$O(n)$	$O(n)$	$O(n)$	$O(n)$
edges()	$O(m)$	$O(m)$	$O(m)$	$O(m)$
getEdge(u, v)	$O(m)$	$O(\min(d_u, d_v))$	$O(1)$ exp.	$O(1)$
outDegree(v) inDegree(v)	$O(m)$	$O(1)$	$O(1)$	$O(n)$
outgoingEdges(v) incomingEdges(v)	$O(m)$	$O(d_v)$	$O(d_v)$	$O(n)$
insertVertex(x)	$O(1)$	$O(1)$	$O(1)$	$O(n^2)$
removeVertex(v)	$O(m)$	$O(d_v)$	$O(d_v)$	$O(n^2)$
insertEdge(u, v, x)	$O(1)$	$O(1)$	$O(1)$ exp.	$O(1)$
remove Edge(e)	$O(1)$	$O(1)$	$O(1)$ exp.	$O(1)$

Graph Algorithms

Data Structures for Graphs

- Edge list




- V is a list of vertices and E is a list of edges. Both can be implemented using doubly linked lists.

Graph Algorithms

Data Structures for Graphs

- Adjacency list

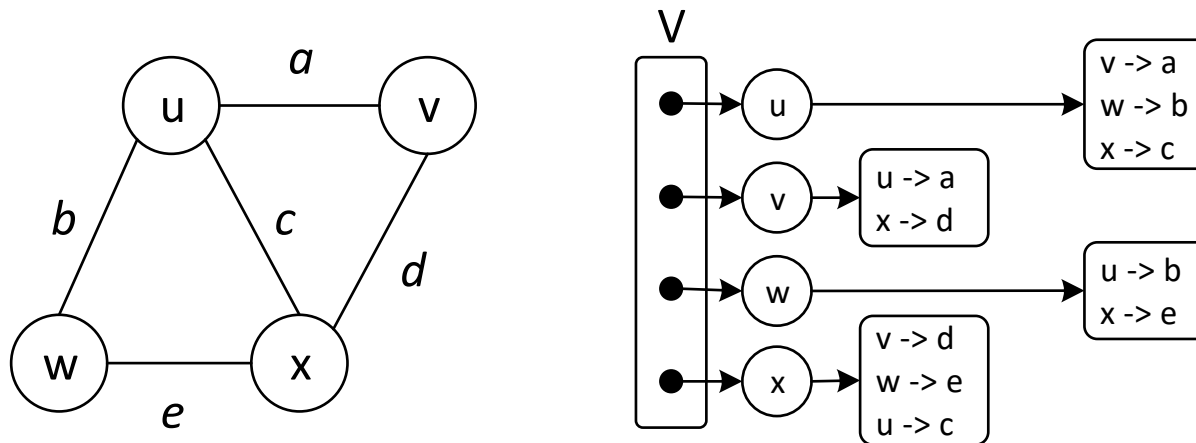


- V is a list of vertices.
- Each vertex v has a reference to a separate collection of edges that are incident to v .
- The collection is called *incidence collection*.

Graph Algorithms

Data Structures for Graphs

- Adjacency map




- Incidence collection of V is implemented as a map.
- Suppose edge $a = (u, v)$ is in the incidence collection of u . Then, $\langle v, a \rangle$ pair is stored in the map, where v is a key and a is the corresponding value.

Graph Algorithms

Data Structures for Graphs

- Adjacency matrix



- $n \times n$ matrix.
- Vertices are encoded to integers and these integers are used as indexes.
- The entry corresponding to vertices u and v stores an edge (u, v) .

Graph Algorithms

Graph Traversals

- A *graph traversal* is a systematic procedure for visiting (and processing) all vertices in the graph.
- We say a traversal is efficient if its running time is proportional to the number of vertices and edges in the graph.
- Applications (for directed graph):
 - Find a direct path from vertex u to vertex v .
 - Find all vertices of G that are reachable from a given vertex s .
 - Determine whether G is acyclic.
 - Determine whether G is strongly connected.

Graph Algorithms

Graph Traversals

- Applications (for undirected graph):
 - Find a path from vertex u to vertex v .
 - Given a start vertex s , find a path with the minimum number of edges from s to every other vertex.
 - Test whether G is connected.
 - Find a spanning tree of G .
 - Identify a cycle in G .
- Will discuss *depth-first search (DFS)* and *bread-first search (BFS)*.

Graph Algorithms

DFS

- Pseudocode

Algorithm DFS (G, u)

Input: A graph G and a vertex u of G

Output: A collection of vertices reachable from u , with
their discovery edges

Mark u as visited

for each of u 's outgoing edges, $e = (u, v)$ do

 if v has not been visited then

 Record edge e as the discovery edge for vertex v

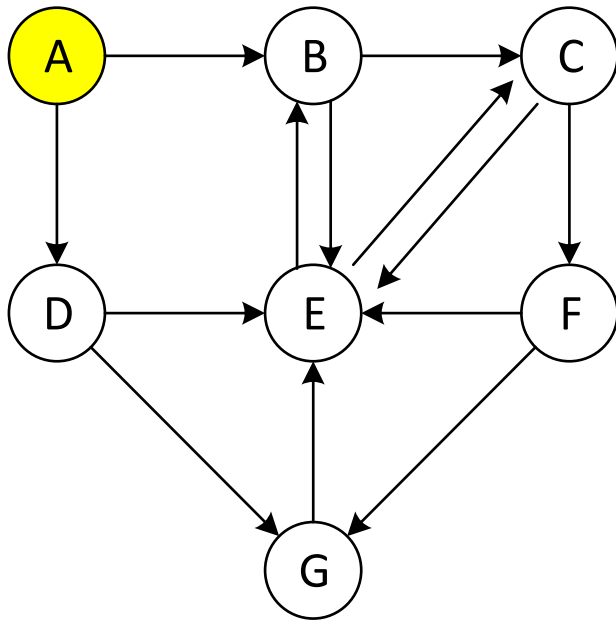
 Recursively call DFS(G, v)

Graph Algorithms

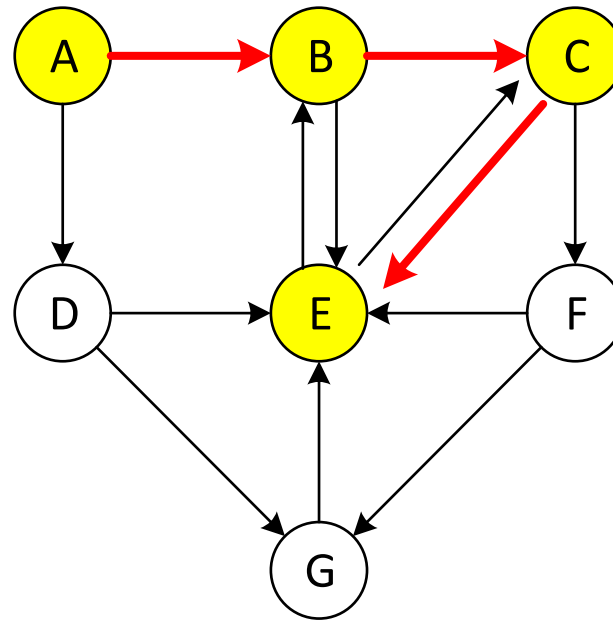
DFS

- Illustration (on a directed graph)

A directed graph
Start at vertex A



A -> B -> C -> E
Backtrack to C



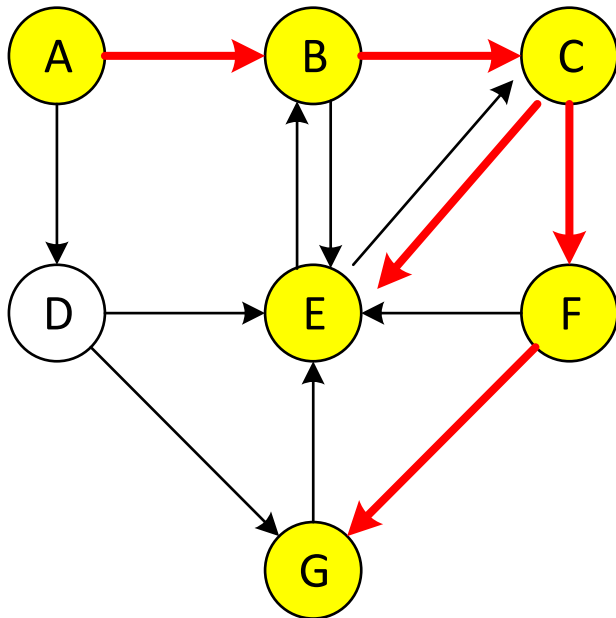
Graph Algorithms

DFS

- Illustration (on a directed graph)


C -> F -> G

Backtrack to F -> C -> B -> A



A -> D

Finished



Graph Algorithms

DFS

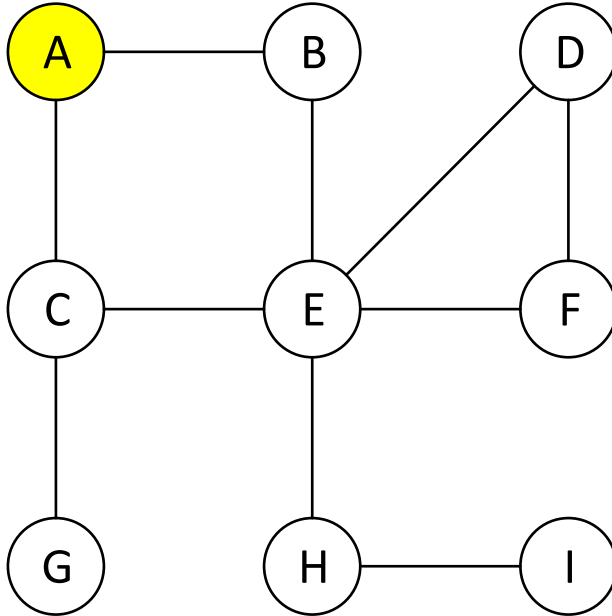
- Illustration (on a directed graph)
 - Classification of edges:
 - *Back edges*: A back edge connects a vertex to its ancestor in the *DFS* tree. They are labeled *B*.
 - *Forward edges*: A forward edge connects a vertex to its descendant in the *DFS* tree. They are labeled *F*.
 - *Cross edge*: A cross edge connects a vertex to a vertex that is neither its ancestor nor its descendant. They are labeled *C*.

Graph Algorithms

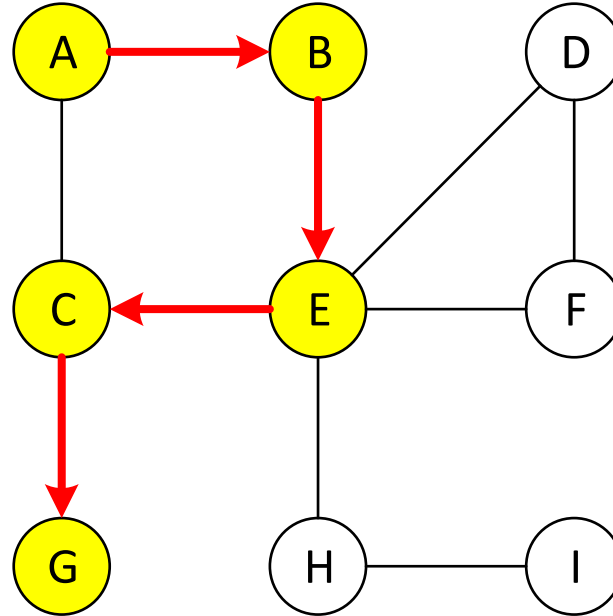
DFS

- Illustration (on an undirected graph)

An undirected graph
Start at vertex A



A -> B -> E -> C -> G
Backtrack to C -> E



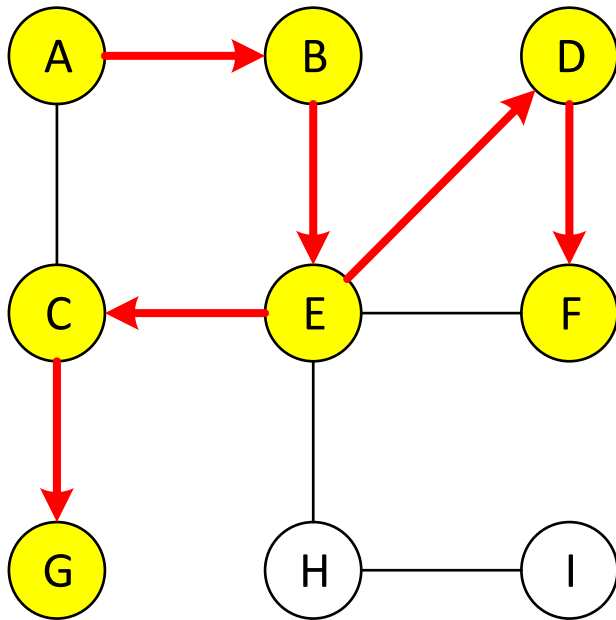
Graph Algorithms

DFS

- Illustration (on an undirected graph)

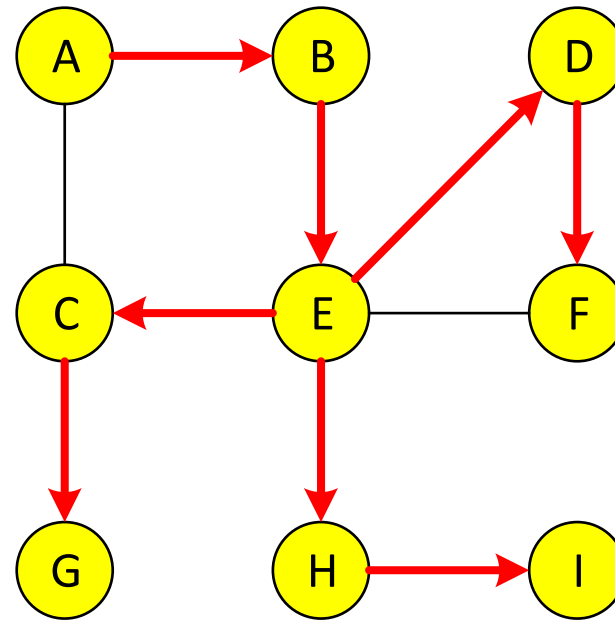
E -> D -> F

Backtrack to D -> E



E -> H -> I

Finished



Graph Algorithms

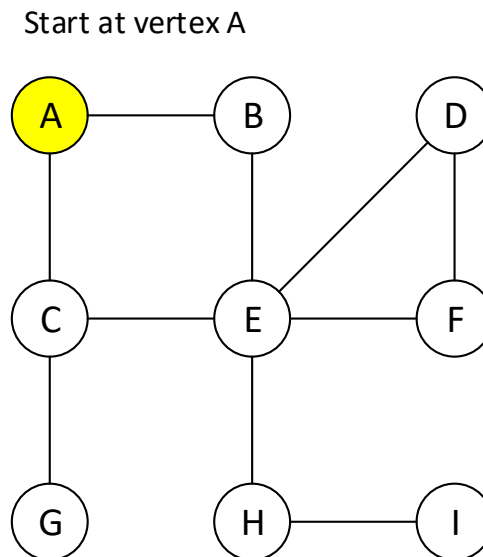
DFS

- DFS properties:
 - A *DFS* on an undirected graph G starting at a vertex s visits all vertices in the connected component of s , and the discovery edges form a *spanning tree* of the connected component of s .
 - A *DFS* on a directed graph G starting at a vertex s visits all vertices reachable from s , and the *DFS* tree contains the directed paths from s to every vertex reachable from s .
- Running time: $O(n_s + m_s)$, here n_s is the number of vertices reachable from s and m_s is the number of edges that are incident to those vertices

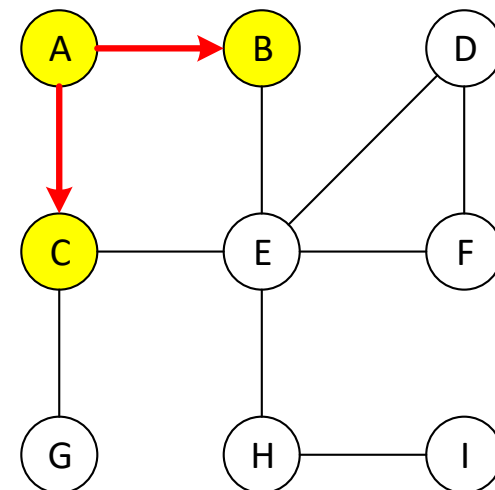
Graph Algorithms

BFS

- Outline
 - Start at the starting vertex s
 - Visit all vertices that are “one-edge away” from s
 - Visit all vertices that are “two-edge away” from s
 - and so on.
- Illustration



Explore vertices that are one-edge away from A.




Graph Algorithms

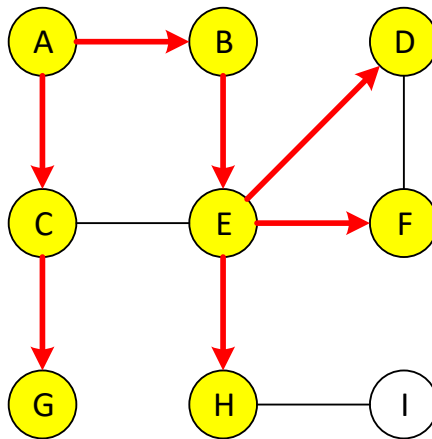
BFS

- Illustration (continued)

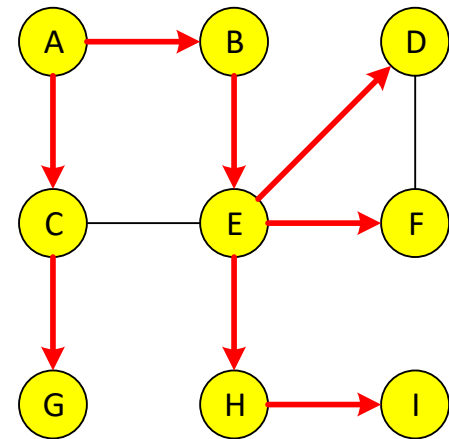
Explore vertices that are two-edge away from A.



Explore vertices that are three-edge away from A.



Explore vertices that are four-edge away from A. Finished



Graph Algorithms


BFS

- BFS properties:
 - The traversal visits all vertices reachable from s .
 - For each vertex v at level i , the path in the *BFS* tree from s to v has i edges, and any other path from s to v in G has at least i edges.
 - If (u, v) is an edge that is not in the *BFS* tree, the level number of v is at most 1 greater than the level number of u .
- Running time: $O(n + m)$

Graph Algorithms

Weighted Graph

- Each edge e is associated with a numeric label called *weight*, denoted $w(e)$.
- Example



Graph Algorithms

Shortest Paths

- Let G be a weighted graph.
 - The *length* of a path P is the sum of the weights of all edges on P . Let $P = \langle (v_0, v_1), (v_1, v_2), \dots, (v_{k-1}, v_k) \rangle$. Then, the length of P , denoted $w(P)$, is defined as:

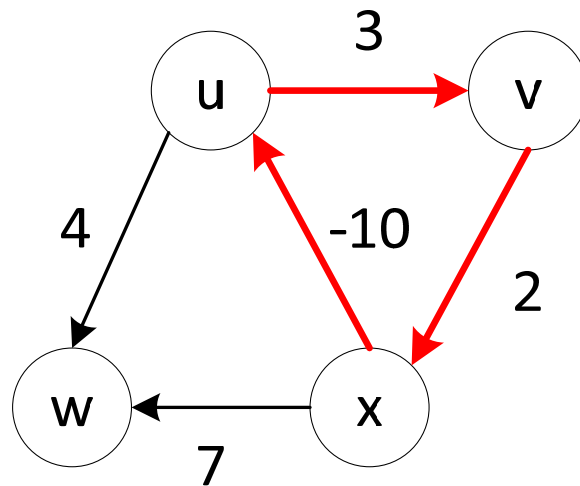
$$w(P) = \sum_{i=0}^{k-1} w(v_i, v_{i+1})$$

- The *distance* from a vertex u to a vertex v in G , $d(u, v)$, is the length of a minimum-length path from u to v , if such path exists. The minimum-length path is referred to as *shortest path*.
 - $d(u, v) = \infty$, if there is no path from u to v in G .

Graph Algorithms

Shortest Paths

- Weights can be negative numbers. Then, a graph may have a *negative-weight cycle*:



- If a graph has a negative-weight cycle, a shortest path is not well defined.

Graph Algorithms

Dijkstra's Algorithm

- A well-known single-source shortest path algorithm on a directed or undirected graph G without negative weights.
- Finds shortest paths from a source vertex to every other vertex in G .
- A greedy algorithm.
- Edge relaxation
 - $D[v]$ is the length of the best path from s to v we have found so far.
 - Initially $D[s] = 0$ and $D[v] = \infty$ for all other vertexes.
 - During the execution of the algorithm, $D[v]$ is updated iteratively and becomes a shortest-path length from s to v .

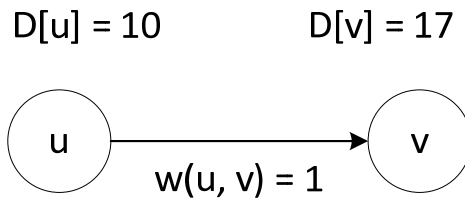
Graph Algorithms

Dijkstra's Algorithm

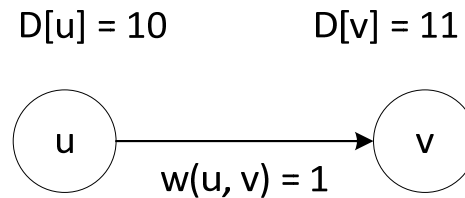
- Edge relaxation (continued)

if $D[u] + w(u, v) < D[v]$ then
 $D[v] = D[u] + w(u, v)$

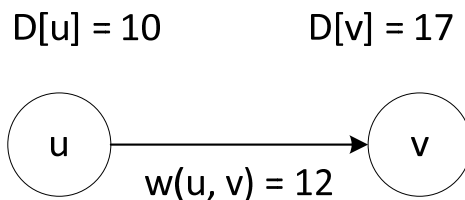
before relaxation



after relaxation



not relaxed



Graph Algorithms

Dijkstra's Algorithm

- Pseudocode

Algorithm ShortestPath(G, s):

Input: A directed or undirected graph G with nonnegative weights, and a distinguished vertex s of G

Output: The length of a shortest path from s to v for every vertex v of G

Initialize $D[s] = 0$ and $D[v] = \infty$ for each vertex $v \neq s$

Let a priority queue Q contains all vertices of G using D labels as keys

while Q is not empty do

$u = Q.\text{removeMin}()$ // vertex with the smallest $D[u]$ is pulled into “cloud”

 for each edge (u, v) such that v is in Q do

 // perform relaxation

 if $D[u] + w(u, v) < D[v]$ then

$D[v] = D[u] + w(u, v)$

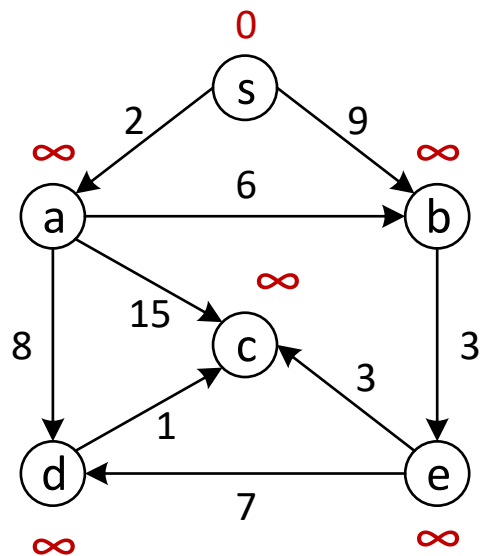
 Change the key of vertex v in Q to $D[v]$

return the label $D[v]$ of each vertex v

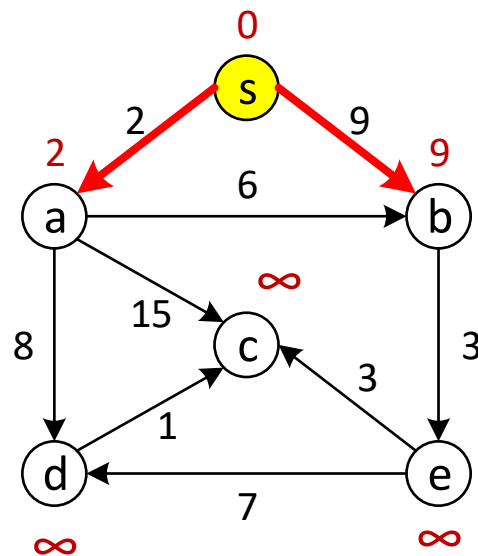
Graph Algorithms

Dijkstra's Algorithm

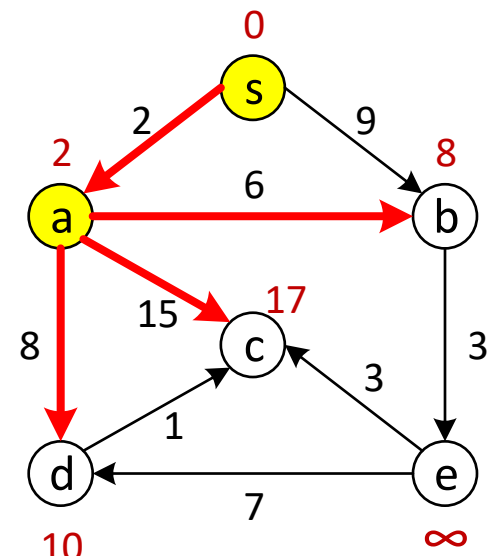
- Illustration



(a) Initially, all vertices are in Q, C is empty, $D[s] = 0$, $D[v] = \infty$ for all other vertices.



(b) s comes into C, edges (s, a) and (s, b) are relaxed.

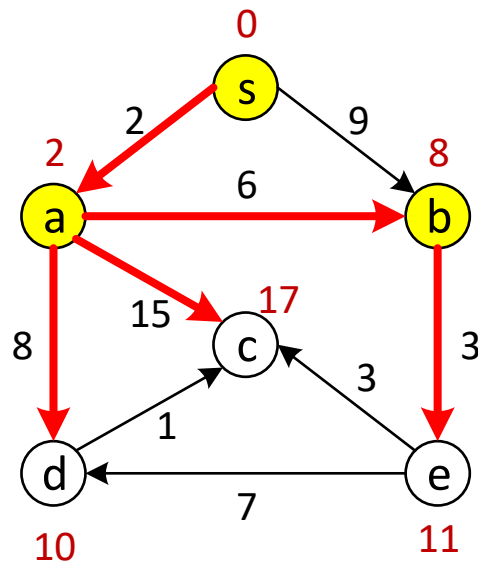


(c) a comes into C, edges (a, b), (a, c), and (a, d) are relaxed.

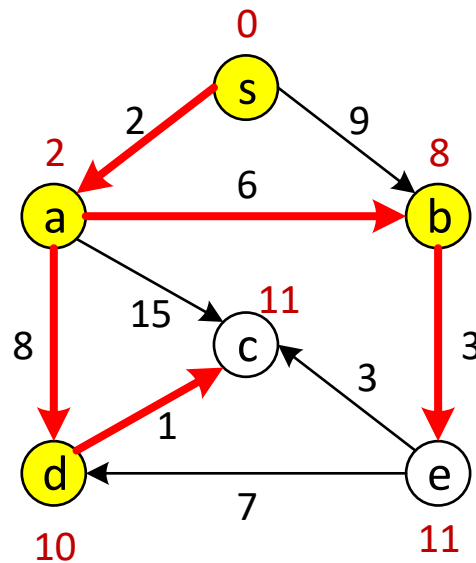
Graph Algorithms

Dijkstra's Algorithm

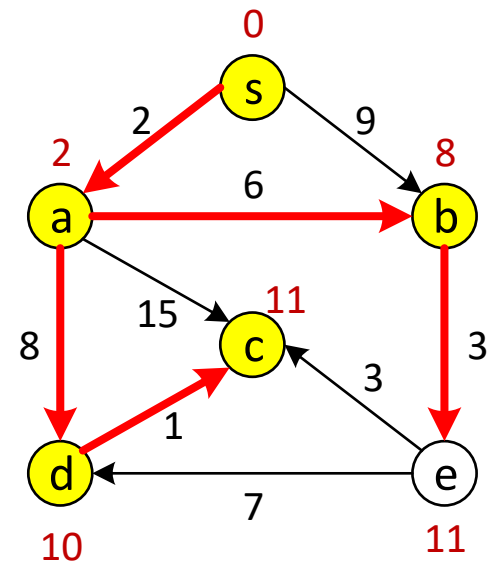
- Illustration (continued)



(d) b comes into C, edge (b, e) is relaxed.



(e) d comes into C, edge (d, c) is relaxed.

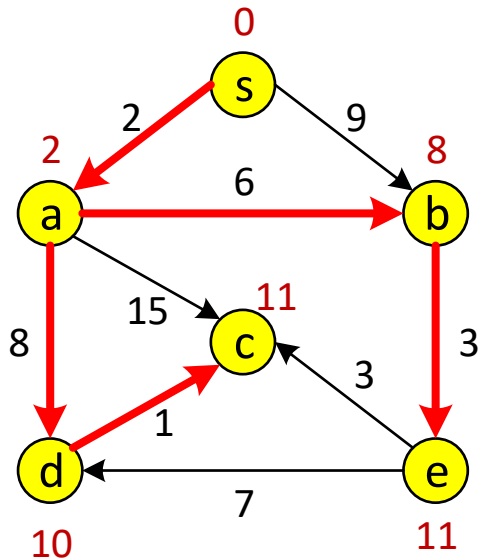


(f) c comes into C. No edge relaxation needed.

Graph Algorithms

Dijkstra's Algorithm

- Illustration (continued)



Running time: $O((n + m) \log n)$

(g) e comes into C. No edge relaxation needed. Finished.

Graph Algorithms

Minimum Spanning Trees

- Given a tree T in an undirected, weighted graph G , the weight of T , $w(T)$, is defined as follows:

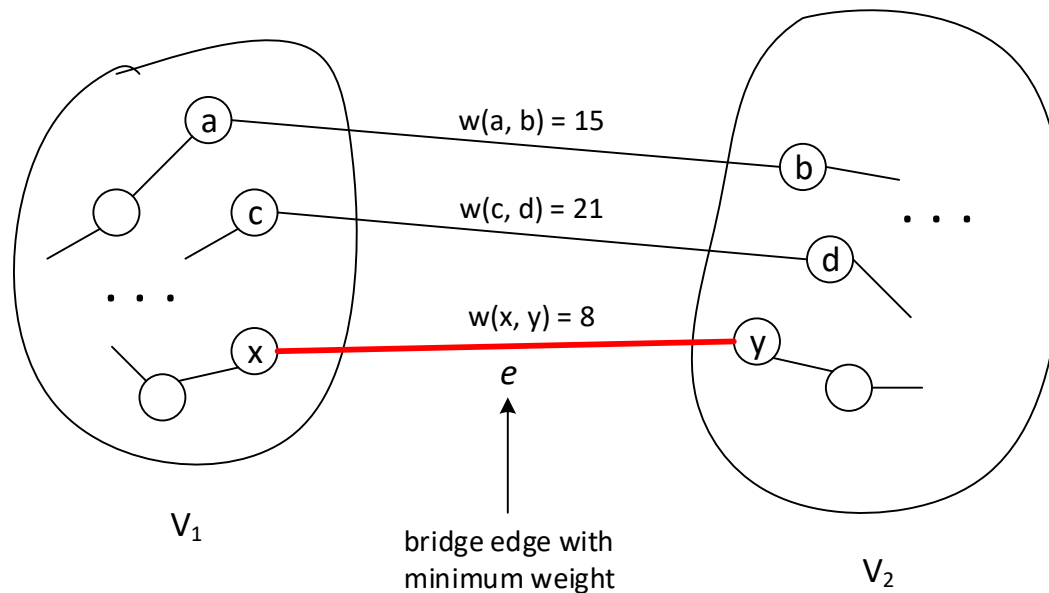
$$w(T) = \sum_{(u,v) \text{ in } T} w(u,v)$$

- A *minimum spanning tree* of an undirected, weighted graph G is a spanning tree with the minimum weight.
- Minimum spanning tree problem: Find such a tree in G .
- Will discuss two algorithms, Prim-Jarnik algorithm and Kruskal's algorithm, both of which are greedy algorithms.
- We assume that a graph G is undirected, weighted, connected, and simple.

Graph Algorithms

Minimum Spanning Trees

- Bridge edge and minimum-weight (bridge) edge
- Suppose G is partitioned into mutually exclusive V_1 and V_2 .
- Bridge edge: one end in V_1 and the other in V_2 .
- Minimum-weight edge: a bridge with the smallest weight



Graph Algorithms

Prim-Jarnik Algorithm

- Outline
 - Begins at some “root” vertex s .
 - Keeps a set of vertices C , called “cloud.”
 - Initially, C has only s .
 - In each iteration, we find a minimum-weight edge connecting a vertex u in the cloud of C and a vertex v that is outside the cloud.
 - Then, the vertex v is pulled into C
 - This process is repeated until a spanning tree is formed.

Graph Algorithms

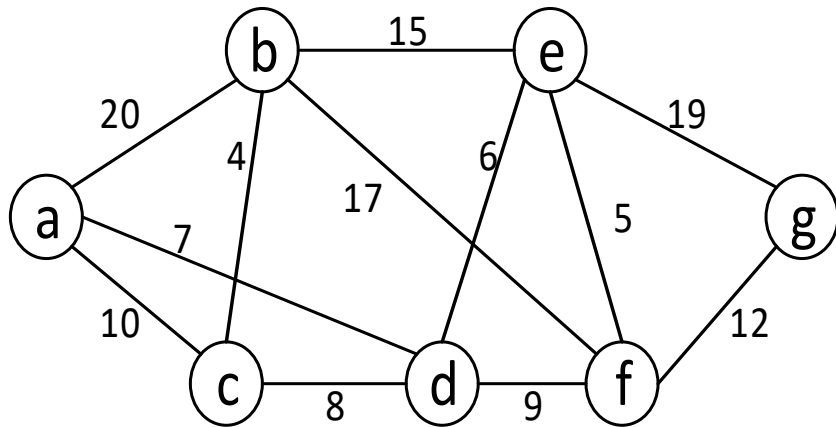
Prim-Jarnik Algorithm

- Outline (continued)
 - Each vertex v has a label $D[v]$, which stores the weight of the minimum observed edge connecting v to the cloud C .
 - Vertices that are not in C are stored in a priority queue, where $D[v]$ is used as a key in the queue.
 - If we choose a vertex in the priority queue with the minimum $D[v]$, then it is a minimum-weight edge.

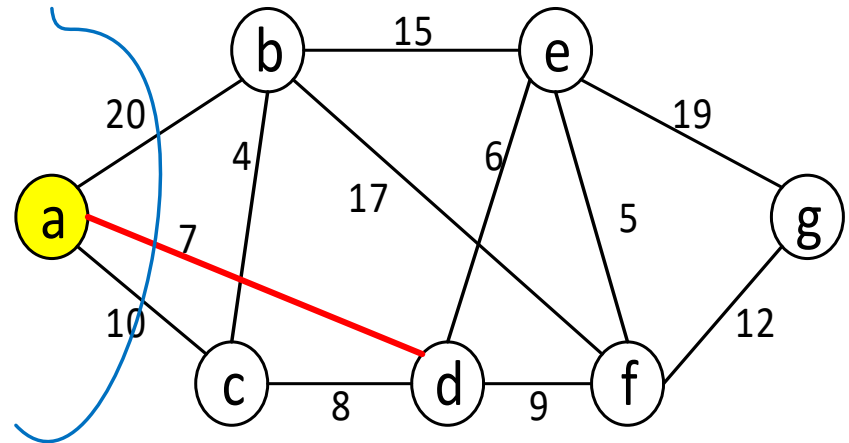
Graph Algorithms

Prim-Jarnik Algorithm

- Illustration
 - Red: min-weight edge;
 - Yellow: in the cloud
 - Blue: cloud boundary



(a) Initial tree. Begin at a




(b) (a, d) is minimum-weight edge


Graph Algorithms

Prim-Jarnik Algorithm

- Illustration (continued)



(c) (d, e) is minimum-weight edge

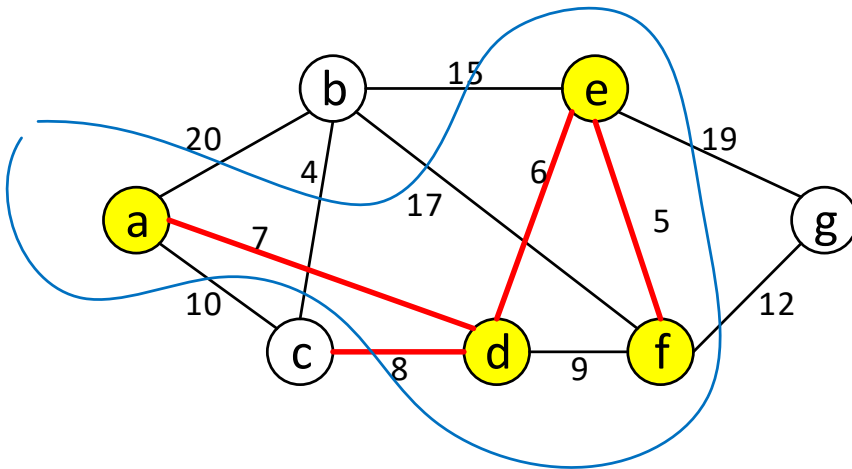


(d) (e, f) is minimum-weight edge

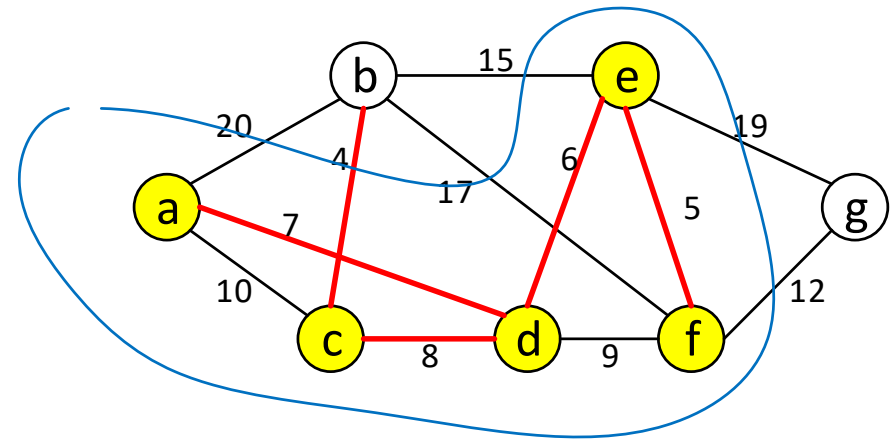
Graph Algorithms

Prim-Jarnik Algorithm

- Illustration (continued)



(e) (c, d) is minimum-weight edge




(f) (b, c) is minimum-weight edge


Graph Algorithms

Prim-Jarnik Algorithm

- Illustration (continued)



(g) (f, g) is minimum-weight edge



(h) Finished. Thick red edges form MST

- Running time: $O((n + m) \log n)$

Graph Algorithms

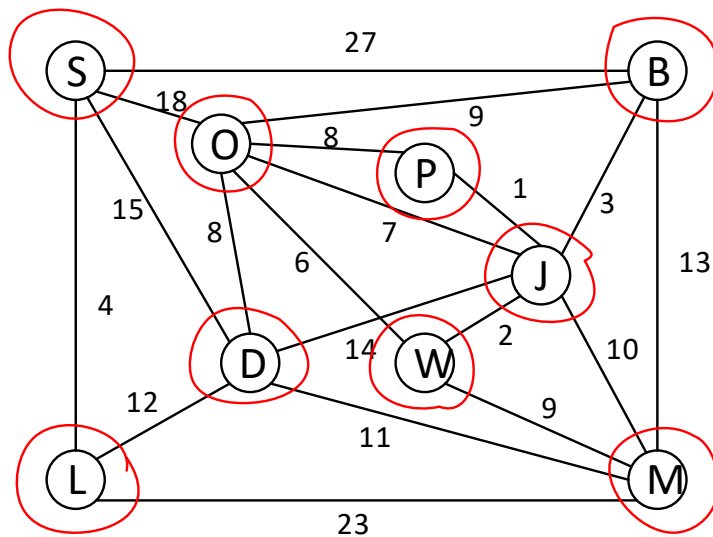
Kruskal's Algorithm

- In the Prim-Jarnik's algorithm, there is always a single tree.
- In the Kruskal's algorithm, there are multiple trees, all of which are eventually merged into an MST.
- Outline: Initially, a spanning tree T is empty and each vertex is a “cluster” on its own.
 - Step 1: Find an edge e with the smallest weight.
 - Step 2: If two endpoints of e belong to different clusters, merge those two clusters.
 - Step 3: Include e in T .
 - Step 4: Stop if all vertices are included by T . Otherwise, return to Step 1 and repeat.

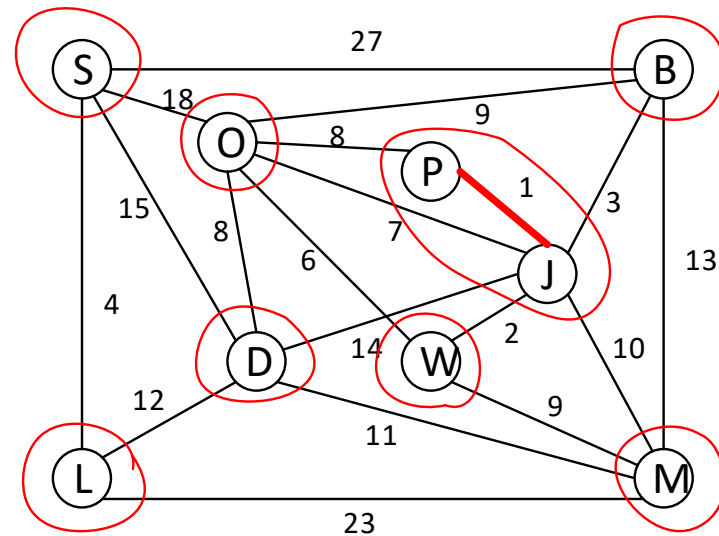
Graph Algorithms

Kruskal's Algorithm

- Pseudocode.
- Illustration



(a) Initial tree. Each vertex is its own cluster. $w(J,P)$ is the smallest.




(b) $w(J,W)$ is the next smallest.

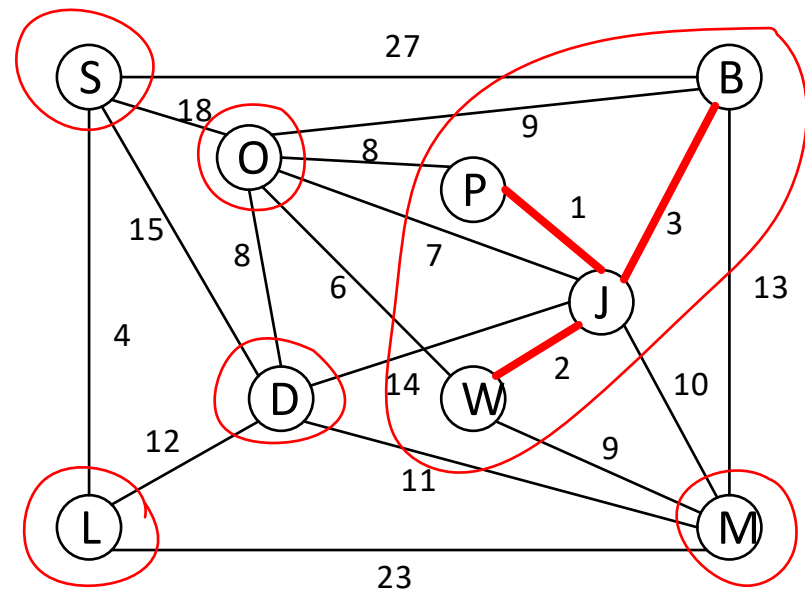
Graph Algorithms

Kruskal's Algorithm

- Illustration (continued)



(c) $w(B,J)$ is the next smallest.




(d) $w(L,S)$ is the next smallest.


Graph Algorithms

Kruskal's Algorithm

- Illustration (continued)



(e) $w(O, W)$ is the next smallest.

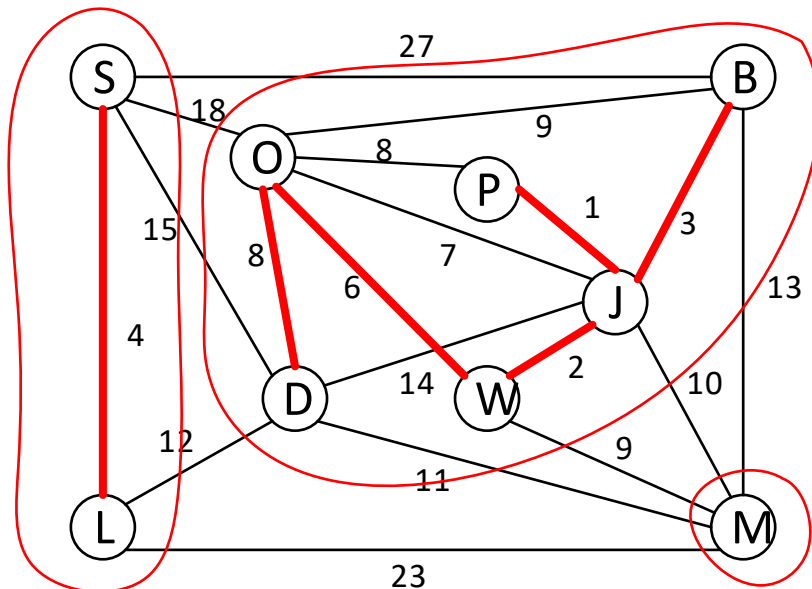


(f) $w(J, O)$ is the next smallest. But, they are in the same cluster. $w(O, P)$ the same. $w(D, O)$ is the next smallest.

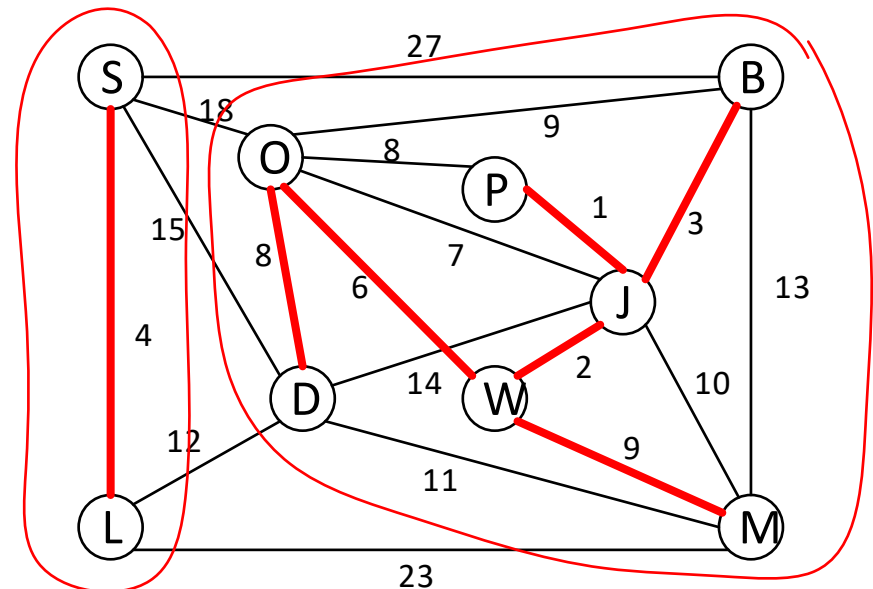
Graph Algorithms

Kruskal's Algorithm

- Illustration (continued)



(g) $w(B,O)$ is the next smallest. But, they are in the same cluster. $w(M,W)$ is the next smallest.

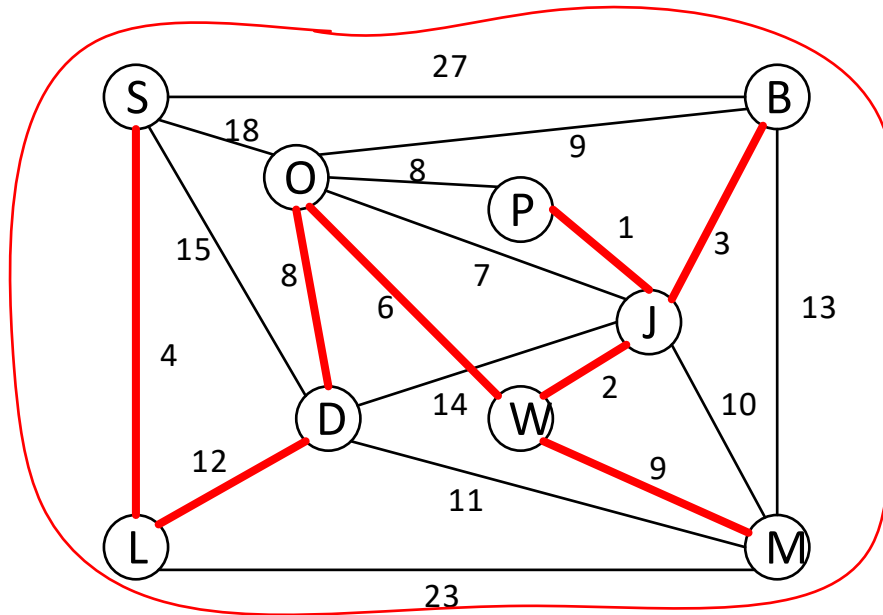


(h) $w(J,M)$ is the next smallest. But, they are in the same cluster. $w(D,M)$ the same. $w(D,L)$ is the next smallest.

Graph Algorithms

Kruskal's Algorithm

- Illustration (continued)



Running time: $O(m \log n)$

(i) Finished. Thick red edges form a minimum spanning tree.

References

- M.T. Goodrich, R. Tamassia, and M.H. Goldwasser, “Data Structures and Algorithms in Java,” Sixth Edition, Wiley, 2014.