

Data Structures and Algorithms


Chapter 11

Binary Search Trees

- Can be used to implement a *sorted map*.
- Each internal position p in a binary search tree stores (k, v) pair.
- Binary search tree is a *proper binary tree* with the following properties:
 - For each internal position p with entry (k, v) pair,
 - Keys stored in the left subtree of p are less than k .
 - Keys stored in the right subtree of p are greater than k .
- Note: In the above definition leaves are “placeholders,” which is shown as small squares in the graph.


Binary Search Trees

- Example (only keys are shown):



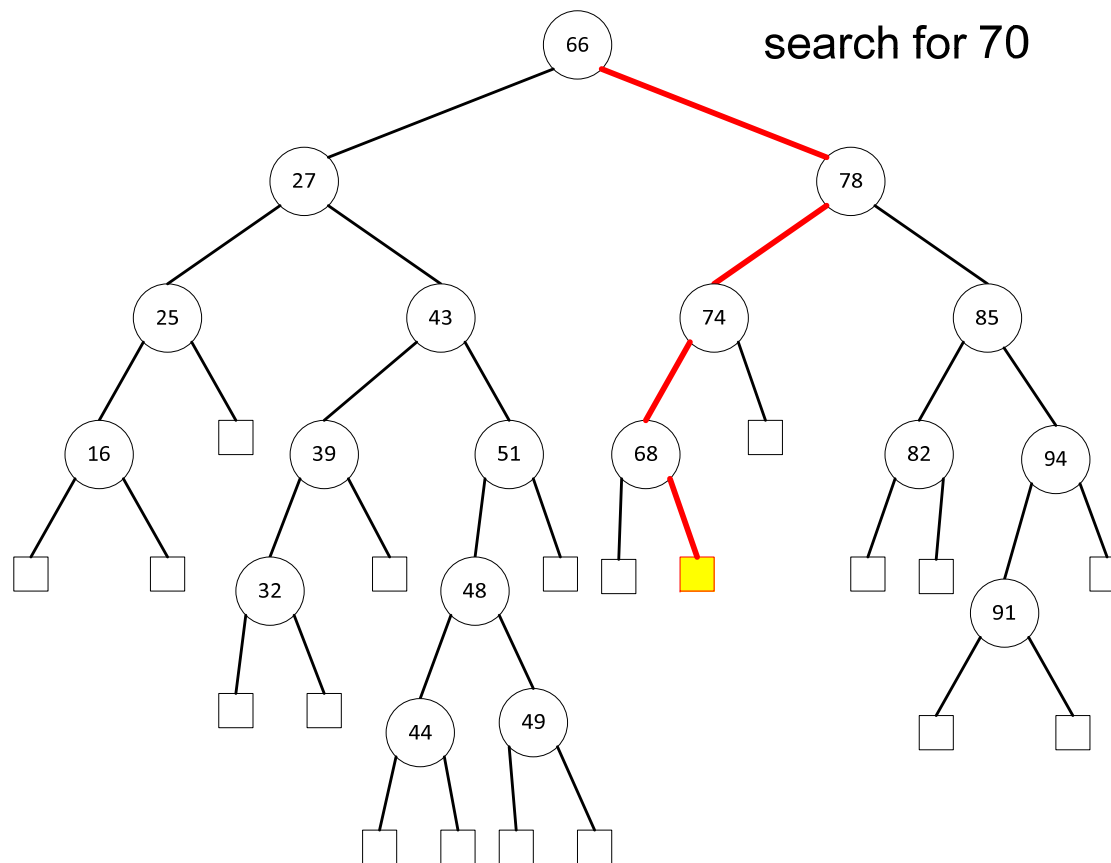
Binary Search Trees

- Search (successful search)



Binary Search Trees

- Search (unsuccessful search)



Binary Search Trees

- Search pseudocode

Algorithm TreeSearch(p , k)

if p is external then // unsuccessful search

return p

else if $k == \text{key}(p)$ // successful search

return p

else if $k < \text{key}(p)$

return TreeSearch(left(p), k) // recurse on left subtree

else

return TreeSearch(right(p), k) // recurse on right subtree

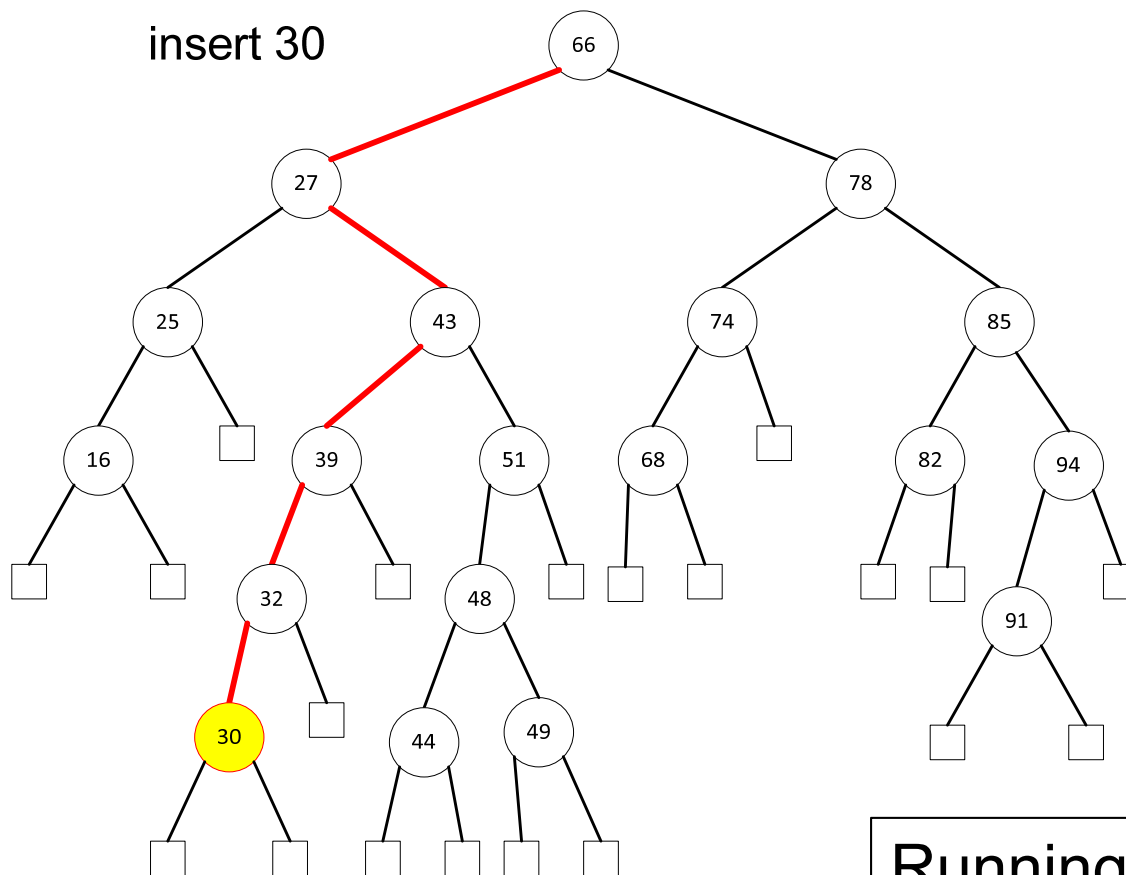
- Running time: $O(h)$

Binary Search Trees

- Inserting an entry with (k, v)
 - Perform a search operation.
 - If an entry with key k is found (i.e., successful search), the existing value is replaced with the new value v .
 - If there is no entry with key k , then we add an entry at the leaf node where the unsuccessful search ended up.


Binary Search Trees

- Insert illustration




Running time: $O(h)$

Binary Search Trees

- Deleting an entry with (k, v) 
 - Slightly more complex
 - Perform search
 - If we reach a leaf node, do nothing
 - If we find the entry at position p
 - Case 1: at most one child of p is an internal node
 - Case 2: p has two children, both of which are internal

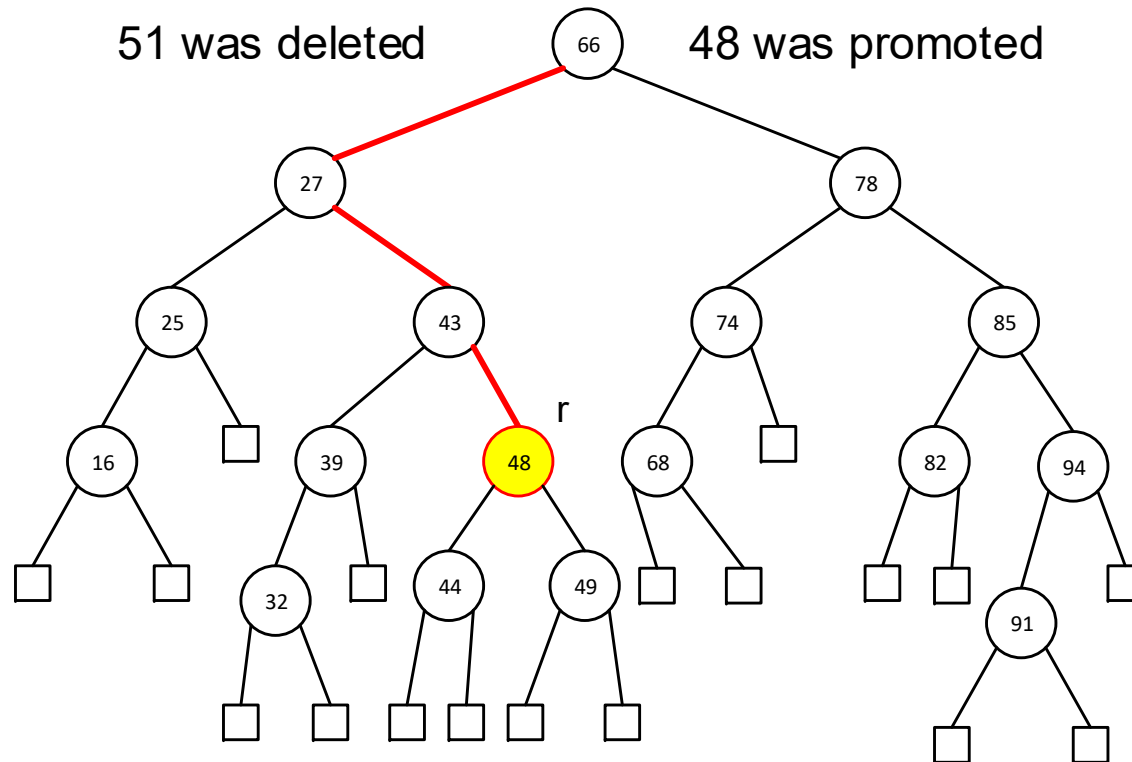
Binary Search Trees

- Deletion Case 1
 - If both children are leaf nodes, then p is replaced with a leaf node.
 - If p has one internal-node child, then that child node replaces p



Binary Search Trees

- Deletion Case 1
 - If p has **one internal-node child** (continued)

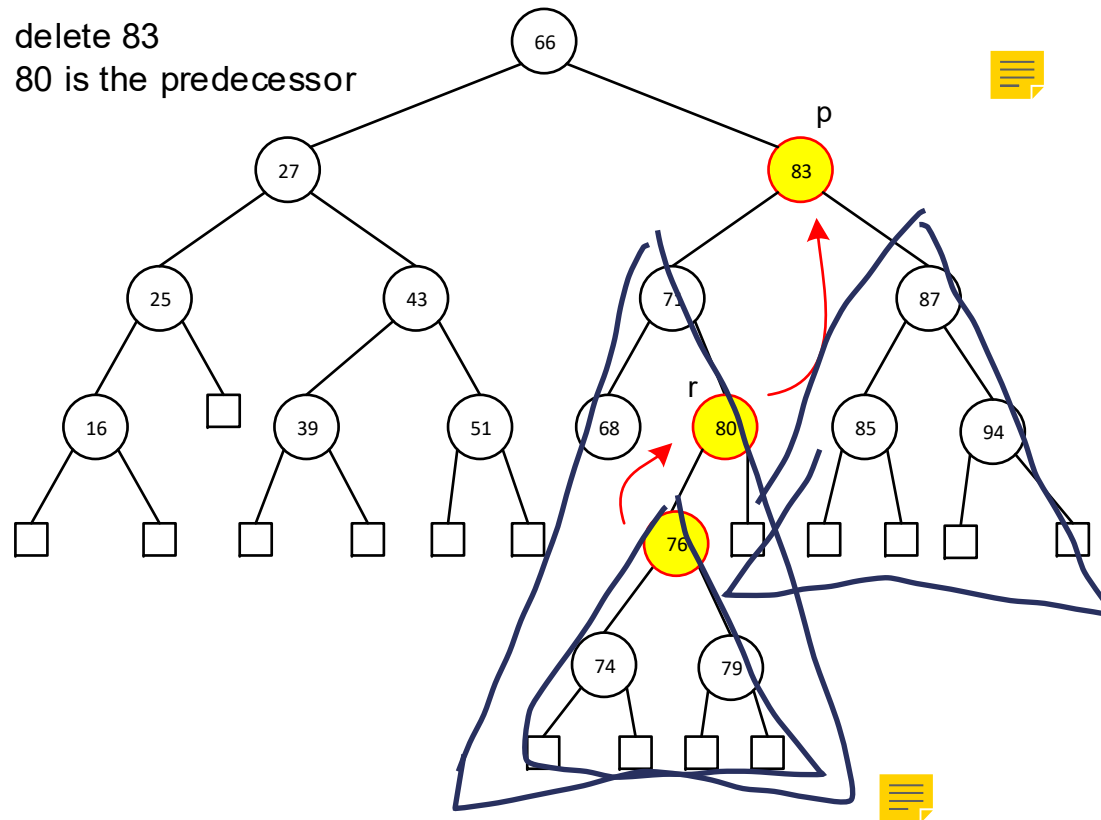


Binary Search Trees

- Deletion Case 2
 - First, we find the node r that has the largest key that is strictly less than p 's key. This node is called the *predecessor* of p in the ordering of keys, which is the rightmost node in p 's left subtree.
 - We let r replace p .
 - Since r is the rightmost node in p 's left subtree, it does not have a right child. It has only a left child.
 - The node r is removed and the subtree rooted at r 's left child is promoted to r 's position.

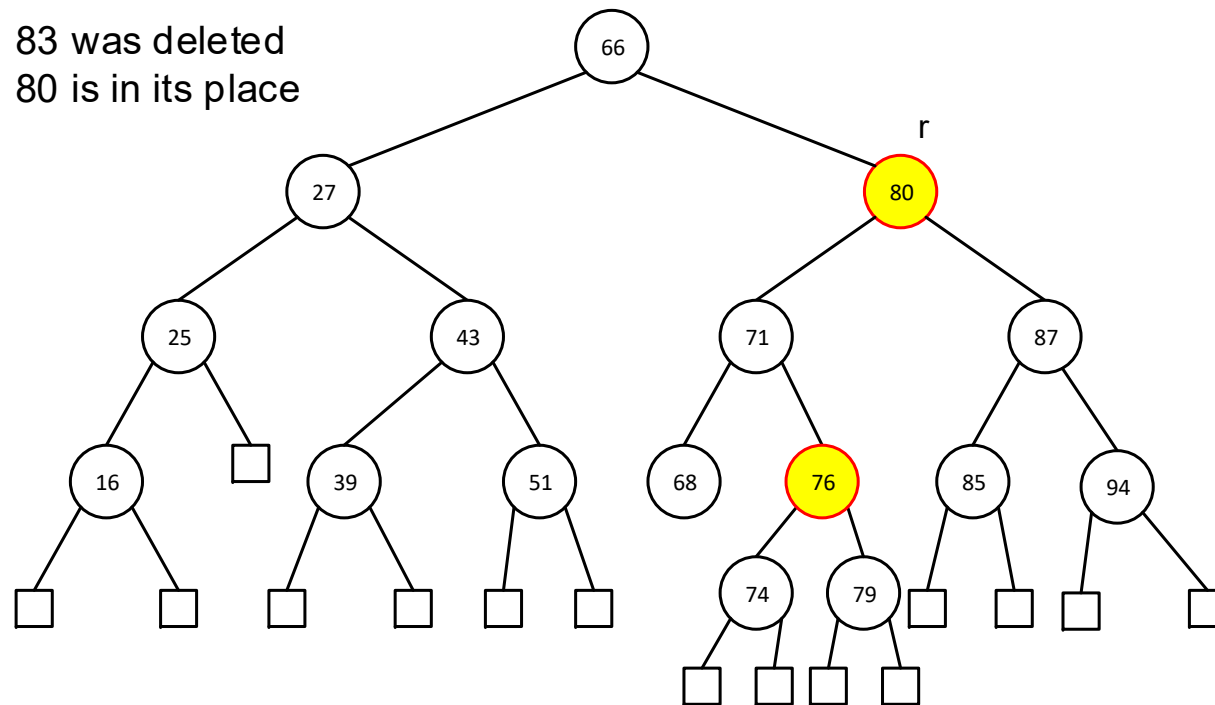
Binary Search Trees

- Deletion Case 2



Binary Search Trees


- Deletion Case 2



- Running time: $O(h)$

Binary Search Trees


- Most binary search tree operations run in $O(h)$.
- In the worst case, a tree is just a linked list. In this case, running times are $O(n)$.



- To guarantee $O(h)$, a tree needs to be balanced.

Balanced Search Trees

- When a binary search tree is unbalanced, it is necessary to *rebalance* the tree.
- Primary operation for rebalancing a binary search tree is *rotation*.



- Can rotate in either direction.
- Binary search tree property is maintained after rotation.

Balanced Search Trees


- A *trinode restructuring* performs a broader rebalancing.
- It involves three positions: x , y , and z
- y is the parent of x and z is the grandparent of x .
- Goal: Restructure the subtree rooted at z to reduce the path length from z to x and its subtrees.
- Use secondary labels, a , b , and c , for the three positions such that a comes before b and b comes before c in an inorder tree traversal of the tree.
- There are four different configurations. These secondary labels allow us to describe the trinode restructuring operations in a uniform way.

Balanced Search Trees

- Outline of the algorithm:
 - (T_1, T_2, T_3, T_4) are left-to-right listing of subtrees of x , y , and z .
 - The subtree rooted at z is replaced with the subtree rooted at b .
 - Make a the left child of b .
 - Make T_1 and T_2 the left and right subtree of a , respectively.
 - Make c the right child of b .
 - Make T_3 and T_4 the left and right subtree of c , respectively.


Balanced Search Trees

- Trinode restructuring: single rotation 1




Balanced Search Trees

- Trinode restructuring: single rotation 2




Balanced Search Trees

- Trinode restructuring: double rotation 1



Balanced Search Trees

- Trinode restructuring: double rotation 2



AVL Trees

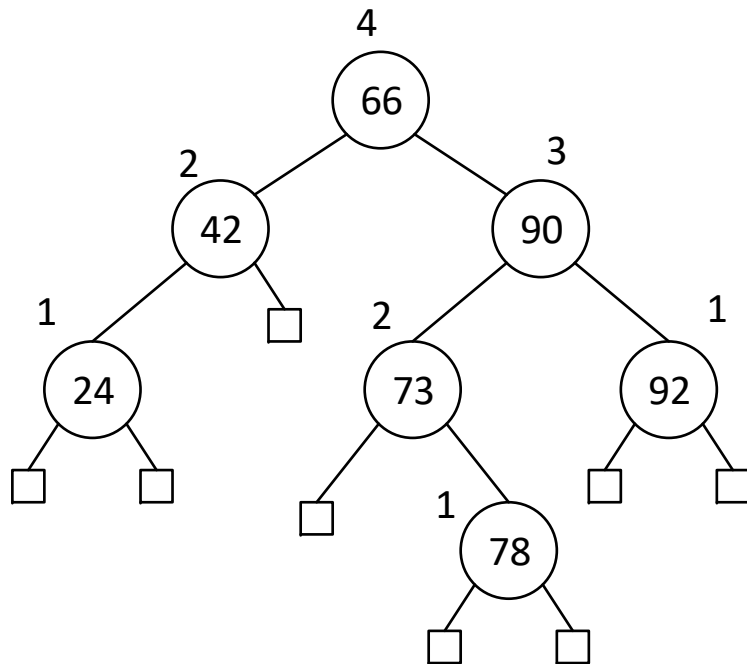
- Recall
 - The *height of a node* is the number of edges on the longest path from that node to a leaf node.
 - The *height of a tree* (or a subtree) is the height of the root of the tree (or a subtree).
 - The height of a leaf node is zero.
- An AVL tree is a binary search tree that satisfies the following *height-balance property*:

For every internal node p of T , the heights of the children of p differ by at most one.

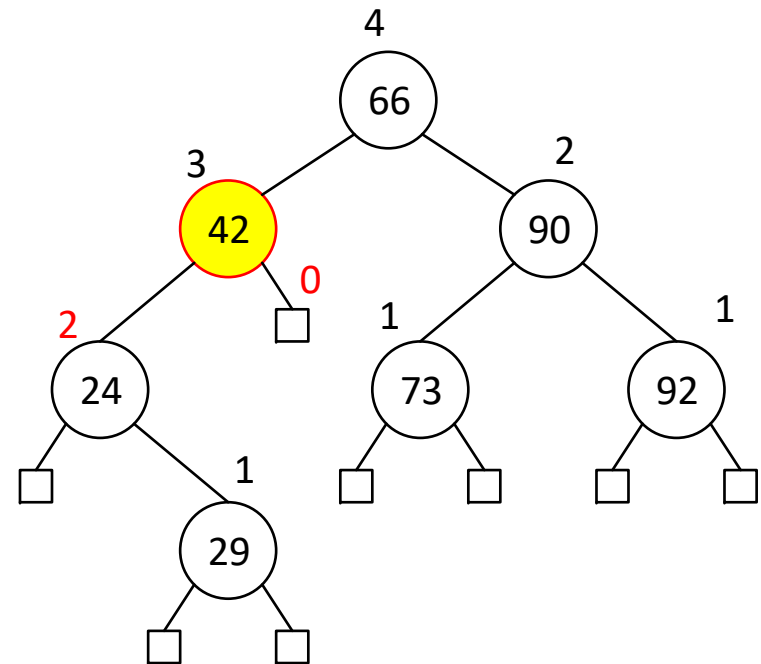
AVL Trees

- AVL tree example:

AVL tree



Not an AVL tree

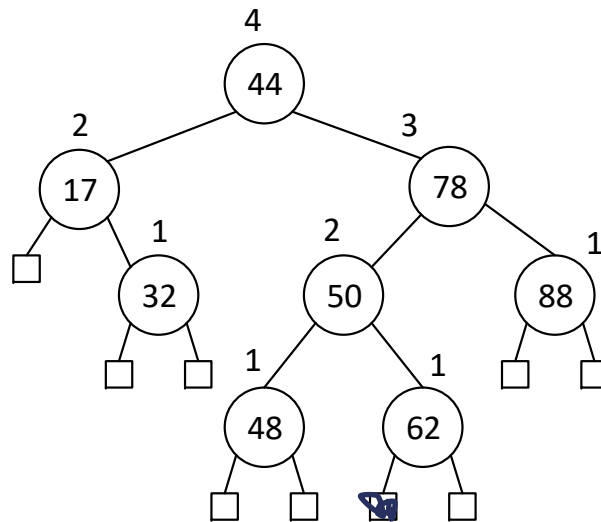


AVL Trees

- Updating an AVL tree
 - A node p in a binary search tree is said to be *balanced* if the heights of p 's children differ by at most one.
 - Otherwise, a node is said to be *unbalanced*.
 - Therefore, every node in an AVL tree is balanced.
 - When we insert a node to an AVL tree or remove a node from an AVL tree, the resulting tree may violate the height-balance property.
 - So, we need to perform *post-processing*.
 - We will discuss only insertion.

AVL Trees

- When a node is inserted, the leaf node p where the new node is inserted becomes an internal node (with the entry of the new node).
- So, ancestors of p may be unbalanced.
- Restructuring is necessary.
- Consider the following tree:




54

54

AVL Trees

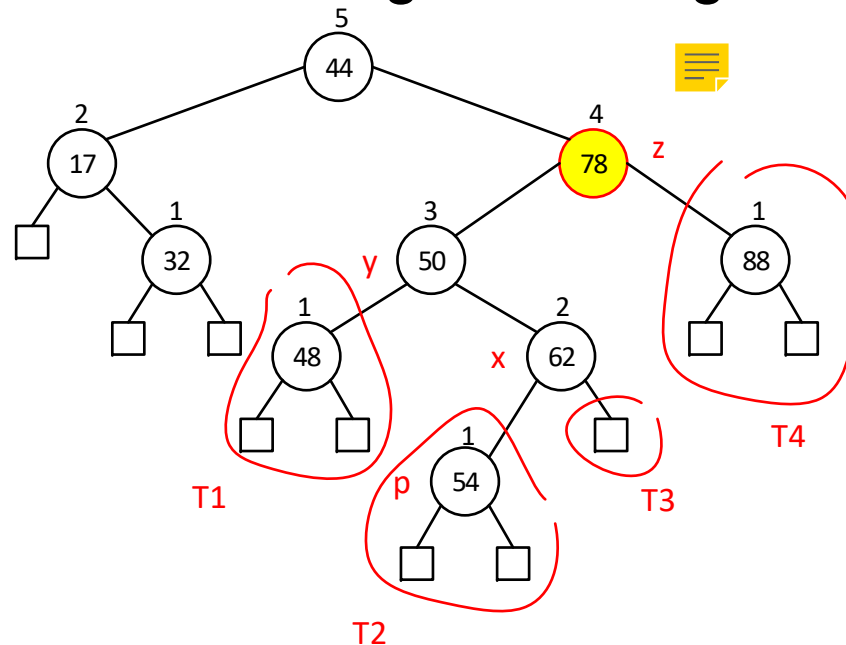
- After inserting 54, the node with 78 is unbalanced

After insertion, before rebalancing



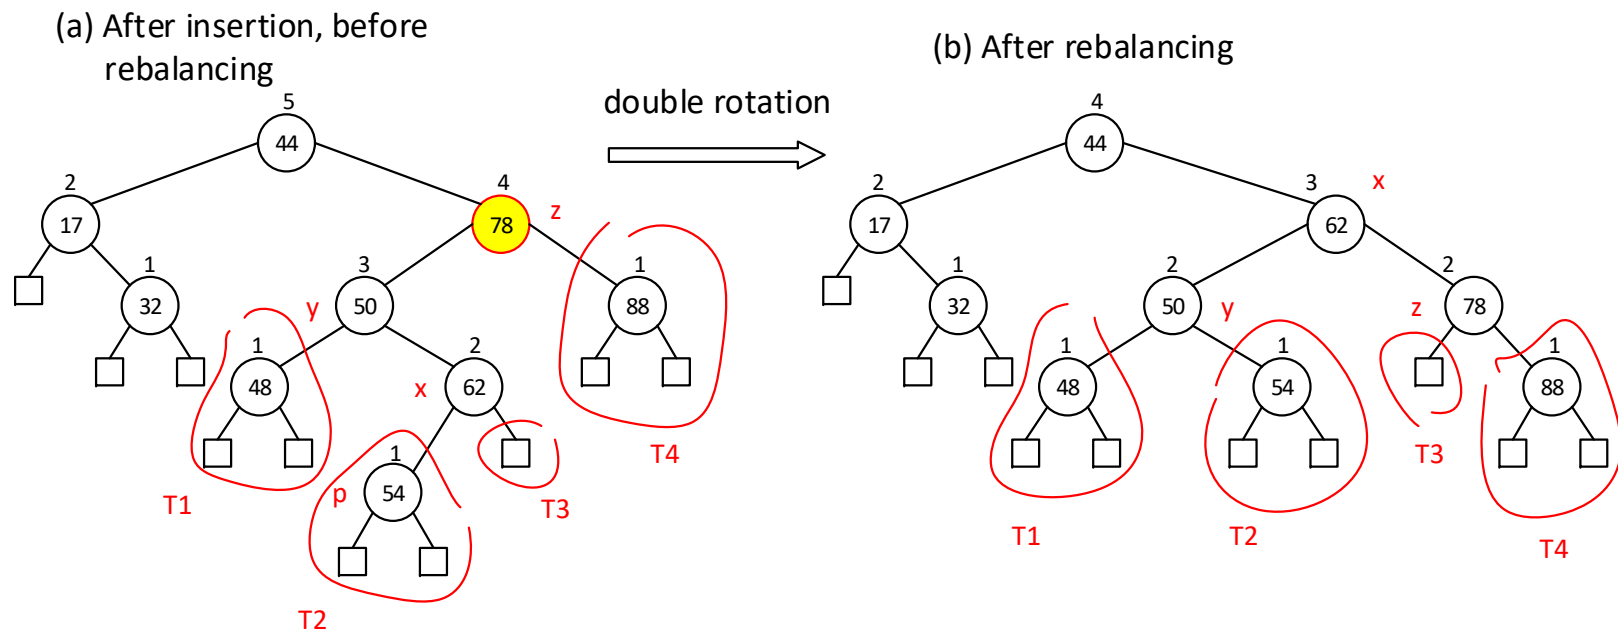
AVL Trees

- Post-processing
 - Search-and-repair strategy
 - Search a node z that is the lowest ancestor of p that is unbalanced.
 - y is z 's child with the greater height
 - x is y 's child with the greater height



AVL Trees

- Perform double rotation to rebalanced the tree



References

- M.T. Goodrich, R. Tamassia, and M.H. Goldwasser, “Data Structures and Algorithms in Java,” Sixth Edition, Wiley, 2014.